

Assignment 3. Reverse-Mode Automatic Differentiation

Due Feb 14 by 11:59pm **Points** 20


Submitting an external tool


Available Jan 31 at 12am - Feb 14 at 11:59pm 15 days

This assignment was locked Feb 14 at 11:59pm.

Due 11:59pm, Feb 14, 2020

Files to submit:

[autodiff.py](#)  (You only need to finish these "TODO" parts, Do Not change the file name)

Please submit to GradeScope Assignment 3. This assignment will be automatically graded. If all your code passes [autodiff_test.py](#)  by run

- `# sudo pip install nose`
- `# nosetests -v autodiff_test.py`

If all your test cases pass, you should see something like this after you submit your code to GradeScope:

Autograder Results

Results

Code

test_add_by_const (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_1 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_2 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_3 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

test_grad_of_grad (test_autodiff.TestAutoDiff) (2.0/2.0)

test_identity (test_autodiff.TestAutoDiff) (2.0/2.0)

test_matmul_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

test_mul_by_const (test_autodiff.TestAutoDiff) (2.0/2.0)

test_mul_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

STUDENT

Jia Zou

AUTOGRADER SCORE

20.0 / 20.0

PASSED TESTS

test_add_by_const (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_1 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_2 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_mul_mix_3 (test_autodiff.TestAutoDiff) (2.0/2.0)

test_add_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

test_grad_of_grad (test_autodiff.TestAutoDiff) (2.0/2.0)


test_identity (test_autodiff.TestAutoDiff) (2.0/2.0)

test_matmul_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

test_mul_by_const (test_autodiff.TestAutoDiff) (2.0/2.0)

test_mul_two_vars (test_autodiff.TestAutoDiff) (2.0/2.0)

Files for testing (This test file include and only include the ten gradescope test cases)

[autodiff_test.py](#)  (You do not need to change this file and you do not need to submit this file)

In this assignment, we would implement reverse-mode auto-diff.

Our code should be able to construct simple expressions, e.g. $y = x_1 * x_2 + x_1$, and evaluate their outputs as well as their gradients (or adjoints: gradients from two output edges), e.g. y , dy/dx_1 and dy/dx_2 .

There are many ways to implement auto-diff, as explained in the slides for Lecture 6. For this assignment, we use the approach of a computation graph and an explicit construction of gradient (adjoint) nodes, similar to what MXNet and Tensorflow do.

Key concepts and data structures that we would need to implement are

- Computation graph and Node
- Operator, e.g. Add, MatMul, Placeholder, Oneslike
- Construction of gradient nodes given forward graph
- Executor

Overview

Here we use a simple example to illustrate the API and data structures of the autodiff module.

Suppose our expression is $y = x_1 * x_2 + x_1$, we first define our variables x_1 and x_2 symbolically,

```
import autodiff as ad


x1 = ad.Variable(name = "x1")
```

```
x2 = ad.Variable(name = "x2")
```

Then, you can define the symbolic expression for y,

```
y = x1 * x2 + x1
```

Now, the computation graph looks like this,



[course/assignment1/blob/master/img/hwk1_graph1.png](https://github.com/dlsys-course/assignment1/blob/master/img/hwk1_graph1.png)

Here, each node is associated with an operator object (we only need a singleton instance for each operator since it is used in an immutable way).

- Node x_1 and x_2 are associated with Placeholder Op.
- Node $(x_1 * x_2)$ is associated with MulOp, and y with AddOp.

With this computation graph, we can evaluate the value of y given any values of x_1 and x_2 : simply walk the graph in a topological order, and for each node, use its associated operator to compute an output value given input values. The evaluation is done in `Executor.run` method.

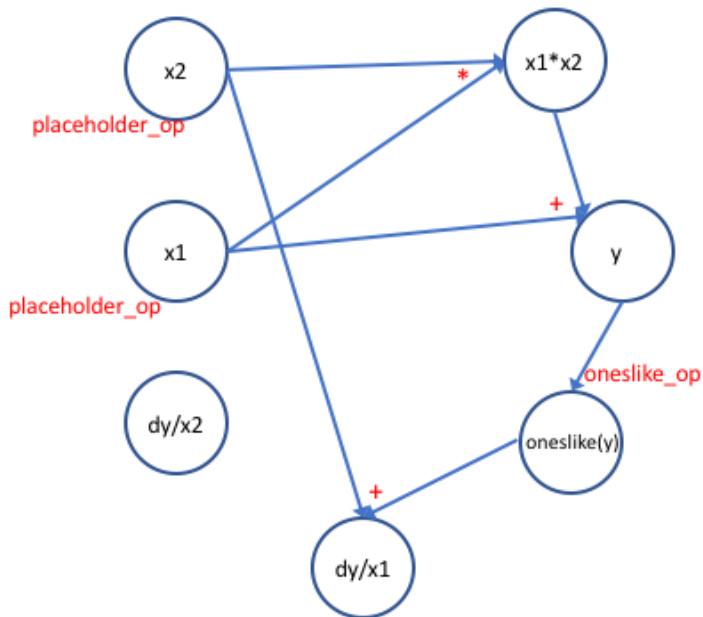
```
executor = ad.Executor([y])
y_val = executor.run(feed_dict = {x1 : x1_val, x2 : x2_val})
```

If we want to evaluate the gradients of y with respect to x_1 and x_2 , as we would often do for loss function wrt parameters in usual machine learning training steps, we need to construct the gradient nodes, `grad_x1` and `grad_x2`.

```
grad_x1, grad_x2 = ad.gradients(y, [x1, x2])
```


According to the reverse-mode autodiff algorithm described in the lecture, we create a gradient node for each node in the existing graph and return those that user are interested in evaluating.

We do this in a reverse topological order, e.g., y , $(x_1 + x_2)$, x_1 , x_2 , as shown in the figures below



[_ \(https://github.com/dlsys-](https://github.com/dlsys-)

[course/assignment1/blob/master/img/hwk1_graph2.png\)](https://github.com/dlsys-course/assignment1/blob/master/img/hwk1_graph2.png)



[_ \(https://github.com/dlsys-](https://github.com/dlsys-)

[course/assignment1/blob/master/img/hwk1_graph3.png\)](https://github.com/dlsys-course/assignment1/blob/master/img/hwk1_graph3.png)

Once we construct the gradients node, and have references to them, we can evaluate the gradients using Executor as before,

```
executor = ad.Executor([y, grad_x1, grad_x2])
y_val, grad_x1_val, grad_x2_val = executor.run(feed_dict = {x1 : x1_val, x2 : x2_val})
```

`grad_x1_val`, `grad_x2_val` now contain the values of dy/dx_1 and dy/dx_2 .

Special Notes

- For simplicity, our implementation expect all variables to have `numpy.ndarray` data type. See tests `feed_dict` usage;

- taking derivative of dy/dx , even though y can be a vector, we are implicitly assuming that we are taking derivative of the `reduce_sum(y)` wrt x . This is the common case for machine learning applications as loss function is scalar or the reduce sum of vectors. Our code skeleton takes care of this by initializing the `dy` as `ones_like(y)` in `Executor.gradients` method.

What you need to do?

- Understand the code skeleton and tests. Fill in implementation wherever marked `"""TODO: Your code here"""`.
- We have 10 tests in `autodiff_test.py`. We would grade you based on those tests.
- Run all tests with `# sudo pip install nose #nosetests -v autodiff_test.py`

Bonus points

Once your code can clear all tests, your `autodiff` module is almost ready to train a logistic regression model. If you are up for a challenge, try

- Implement all missing operators necessary for a logistic regression, e.g. `log`, `reduce_sum`.
- Write a simple training loop that updates parameters using gradients computed from `autodiff` module.

Grading Rubrics

- `autodiff_test.test_identity` ... 2 pt
- `autodiff_test.test_add_by_const` ... 2 pt
- `autodiff_test.test_mul_by_const` ... 2 pt
- `autodiff_test.test_add_two_vars` ... 2 pt
- `autodiff_test.test_mul_two_vars` ... 2 pt
- `autodiff_test.test_add_mul_mix_1` ... 2 pt
- `autodiff_test.test_add_mul_mix_2` ... 2 pt
- `autodiff_test.test_add_mul_mix_3` ... 2 pt
- `autodiff_test.test_grad_of_grad` ... 2 pt
- `autodiff_test.test_matmul_two_vars` ... 2 pt
- bonus (training logistic regression) ... 1 bonus pt added to your final grade